

Learning and distinguishing between active versions of prominent decision strategies

Eric Schulz¹, Paula Parpart¹,
Björn Meder², Bradley C. Love¹ & Maarten Speekenbrink¹

¹UCL London, Department of Experimental Psychology

²MPI Berlin, Center for Adaptive Behavior and Cognition

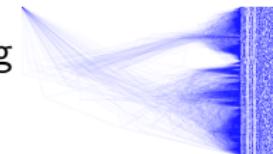
June 18, 2015

If humans are active learners, then...

- 1 ...how can they learn a heuristic that is made up of building blocks on the fly when the goal is to make good decisions?



- 2 ...how can we use that fact to test different decision making strategies in an active learning context?



Overview

1 Heuristics from Approximate Bayesian Computation

Collaborators:  Meder



Speekenbrink

- How can a heuristic emerge from small building blocks?
- Can a ABC method recover and/or outperform TTB?

2 Distinguishing active versions of decision strategies

Collaborators:  Parpart  Speekenbrink  Love

- How to distinguish between active versions of strategies?
- How does the used strategy depend on the environment?

PART 1: Can heuristics emerge/grow adaptively?



- Heuristics are made of smaller building blocks
- Different combinations of blocks produce the heuristic toolbox
- Definition of a heuristic:

"Ignores information to be faster and/or more accurate."

"Exhibits a starting rule, a search rule, and a stopping rule."

PART 1: Problem statement

Problem

Given a number of binary cues C_1, C_2, \dots, C_n –some of which might be invalid– predict a binary outcome y on every trial while at the same time learning a heuristic structure that consists of a starting, a search, and a stopping rule.

Take-The-Best cannot deal with that, because...

- 1 it is not clear how cue validities are learned
- 2 it is not clear how information is ignored
- 3 it cannot make decision and learn (evolve over time)
- 4 it is impossible to compute all validities exactly on the fly

PART 1: Approximate Bayesian Computation

How do you calculate a posterior when both analytical inference and sampling methods are intractable?

Requirements: model \mathcal{M} , prior $\pi(\theta)$, data $\mathcal{D} = (\mathbf{X}, \mathbf{y})$

- 1 Sample a proposal θ^* from your prior $\pi(\theta)$
- 2 Plug the proposal in your model \mathcal{M} to get a model \mathcal{M}^*
- 3 Simulate $\mathbf{y}^* = \mathcal{M}^*(\mathbf{X})$
- 4 Calculate distance $\delta = \sum_i^n |y_i - y_i^*|$
- 5 If $\delta < \epsilon \rightarrow$ accept proposal
else \rightarrow reject proposal

PART 1: Approximately Bayesian Computed TTB

How can TTB emerge from smaller building blocks?

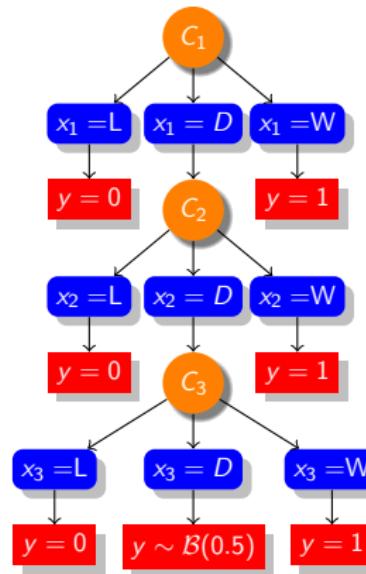
Requirements: urn containing the cues, for each cue: an urn containing the weights

urn has a pre-assigned number of k balls for each class

- 1 Sample a cue c^* from the cue urn without replacement
- 2 Sample a weight w^* for c^* from c^* 's weight urn
- 3 Encounter a situation y_i
- 4 If $y^* == y_i \rightarrow$ put two more balls into w^* and c^*
elseif $y^* == 0 \rightarrow$ go back to 1
elseif $y^* \neq y_i \rightarrow$ reject proposal, start again

PART 1: Recovering TTB-The set up

- 5 variables x_1, x_2, x_3, x_4, x_5 , only the first three are valid
- $p(x_k = W) = p(x_k = D) = p(x_k = L) = \frac{1}{3}$



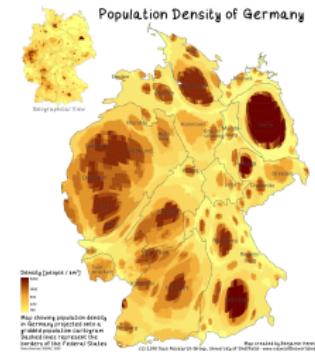
PART 1: Recovering TTB: Results

- Results
 - 1 Can find the most important cues, even with a greedy strategy
 - 2 Achieve an excellent misclassification rate
 - 3 Can show mathematically that TTB-ABC corresponds to classic TTB if full set is sampled set
- Current problems
 - 1 How to avoid local optima (rescaling, jump starts, sampling)?
 - 2 How to speed up computation (annealing, parallelization)?

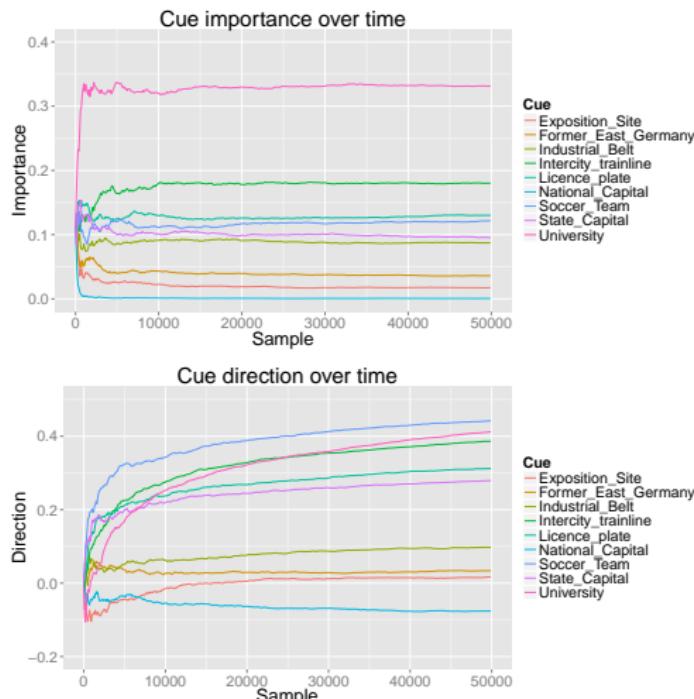
Cue importance over time

PART 1: Empirical Test: City size data

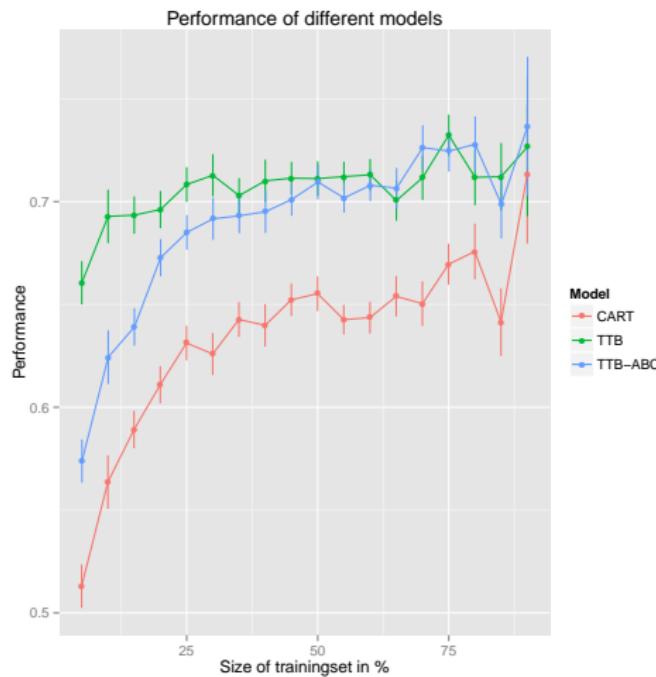
- Most frequently used data sets within the community
- Predict which of two German cities has the larger population size based on 9 cues
- Take-The-Best is known to perform well
- Cues have been pre-selected



PART 1: Empirical Test: Importance and direction



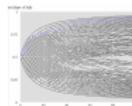
PART 1: Empirical Test: Model comparison



PART 1: Conclusion



TTB can emerge from simple building blocks over time



Relies on simple sampling and approximate computation



Can recover and perform as well as simple decision trees



Speed, global optima, and experimental applications

PART 2: Can we distinguish different active algorithms?



Golf the Volkswagen compact.
Great tax, great performance and great fun.

- Agent has evolved to obey a certain decision algorithm
- Agent learns with the goal to apply that algorithm
- Definition of active learning:

“Stepwise selection of observations in order to learn faster.”

“Active learning leads to a banana shaped learning curve.”

PART 2: Problem statement

Problem

Given a number of binary cues C_1, C_2, \dots, C_n –some of which might be invalid– learn how to predict a binary outcome y by selecting observations as wisely as possible.

Current heuristics cannot deal with that, because...

- 1 it is not clear how cue validities are learned
- 2 it is not clear how information is ignored
- 3 there are no active versions of them
- 4 active learning has never been utilized for model comparison

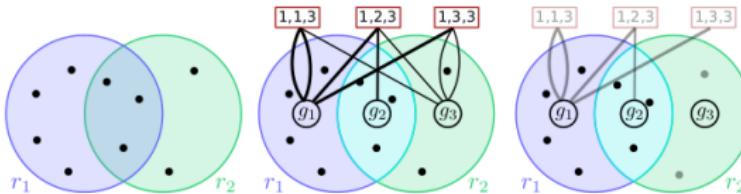
PART 2: Model 1: Active TTB

- **Imagine a scenario without noise:**

Every observation is deterministic and your only goal is, given your hypothesis space over all cue orders,

$\mathcal{H} = \{\text{cue-order}_1, \text{cue-order}_2, \dots, \text{cue-order}_k\}$, to select the observation $s^* = \text{argmax}\{|\mathcal{H}| - |\mathcal{H}|s|\}$

- Unfortunately, observations never come without noise
- We have to find a probabilistic version of the same algorithm



PART 2: Model 1: Active TTB algorithm

- Put a pseudo-count π over all possible cue orders
- If a cue order makes a correct prediction, then $\pi ++$
- Calculate current entropy over all cue orders $S_0 = \sum_i p_i \log p_i$
- For every possible comparison s calculate $p(y = w)$ and $p(y = L)$ by π -weighted sum over all cue orders
- For every comparison s , calculate posterior expected entropy $\mathbb{E}[S|s] = S(y = W) \times p(y = W) + S(y = L) \times p(y = L)$
- Choose $s^* = \operatorname{argmax}\{S_0 - \mathbb{E}[S|s]\}$, that is the observation with the highest expected information gain
- Method works well a priori

PART 2: Model 2: Active Logistic Regression

- Given a Bayesian variant of logistic regression:

$$f(x) = \frac{1}{1 + \exp(-(\beta_0 + \sum_k \beta_k x_k))}$$

- Calculate the current sum of coefficients' uncertainty

$$S = \sum_k \mathbb{V}(\beta_k)$$

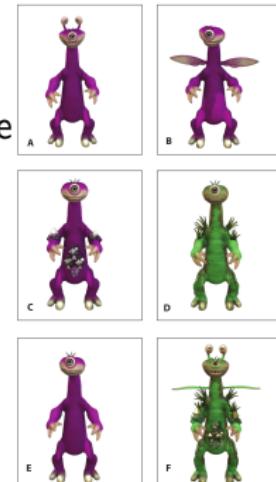
- For every comparison s , calculate $p(y = w|s)$ and $p(y = l|s)$
- For every comparison and outcome calculate $S|s, y = w$ and $S|s, y = l$, that is the reduction of variance
- For every comparison s , calculate posterior expected uncertainty

$$\mathbb{E}[S|s] = S(y = W) \times p(y = W) + S(y = L) \times p(y = L)$$

- Choose $s^* = \operatorname{argmax}\{S_0 - \mathbb{E}[S|s]\}$, that is the observation with the highest expected uncertainty reduction
- Method works well a priori

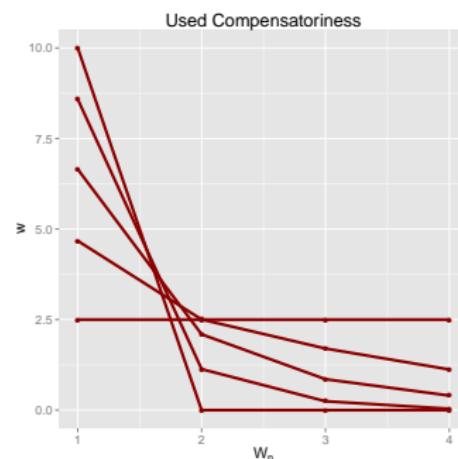
Part 2: Alien Olympics: Design

- 1 Learn how well aliens perform in Olympics
- 2 Aliens vary on 4 different features:
Wings, Camouflage, Diamonds, and Antennæ
- 3 30 learning trials: select 2 out of 4 aliens to compete against each other (\$0.5 reward)
- 4 10 test trials: select 1 out of 2 aliens for your team (bonus dependent on team)
- 5 Environment set up with different levels of compensatoriness

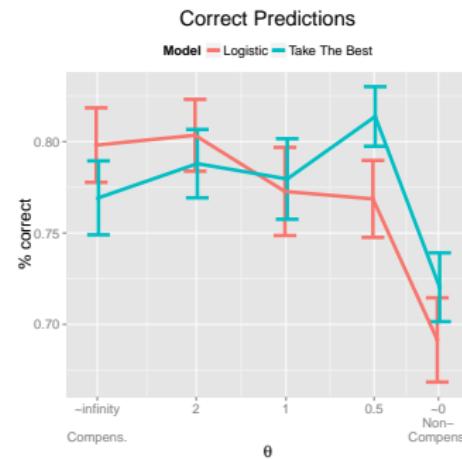
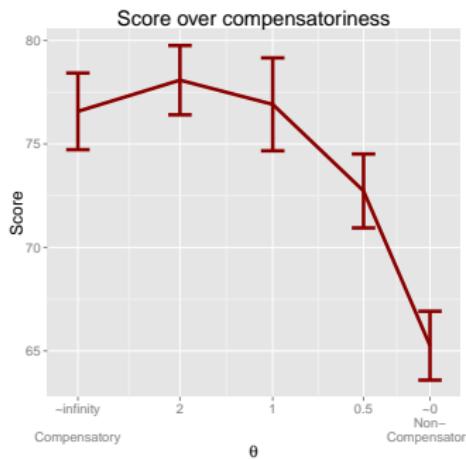


Part 2: Alien Olympics: Results in test set

- 1 Underlying environment obeys logistic regression
- 2 Weights are generated by a stick-breaking process
 $\beta'_k \sim \text{Beta}(1, \theta)$
 Define $\{\beta'_k\}_{k=1}^4$ as:
 $\beta_k = \beta'_k \prod_{i=1}^{k-1} (1 - \beta'_i)$
- 3 Allows to trade-off different levels of compensatoriness

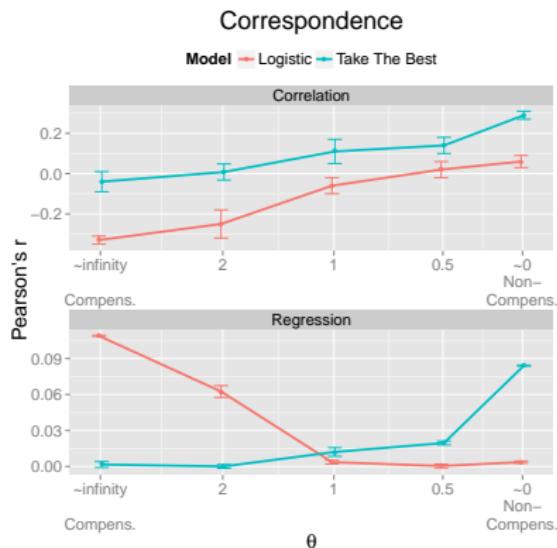
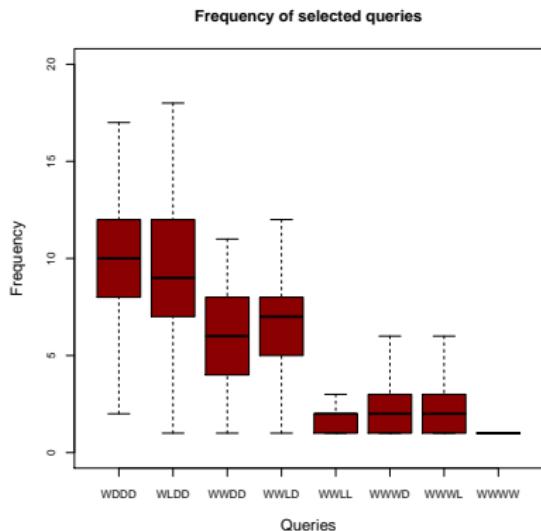


Part 2: Alien Olympics: Results of passive part



- 1 Non-compensatory conditions produce lower average score
- 2 Hard to distinguish between models based on test set alone

Part 2: Alien Olympics: Results of active part

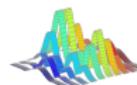


- 1 Participants perform simple but sensible queries
- 2 Can distinguish between active models:
Only active TTB seems to match behavior well

PART 2: Conclusion



Active learning can be used to distinguish different models



It is possible to design active versions of classic models



In a first experiment, active TTB seems to do best



Simpler follow-up, Exploitation scenarios, Non-parametrics

Overall conclusion

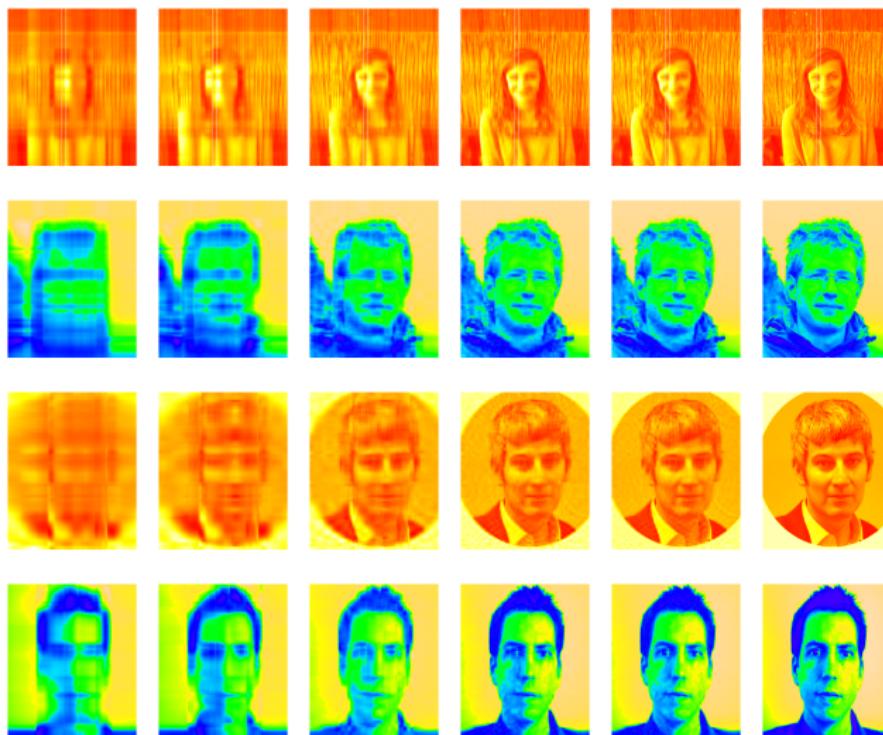
I hope to have convinced you that...

- 1** ...active learning is an exciting tool to add to our methodological repertoire
- 2** ...it is both possible and fruitful to design active learning algorithms based on classic models
- 3** ...if we want to make claims about cognitive strategies, we also have to make claims about how those are acquired

Future steps could involve...

- 1** ...coming up with more active learning models
- 2** ...designing additional active learning algorithms
- 3** ...modeling both exploration and exploitation scenarios

Thanks you for your attention and to my collaborators



Schulz, Parpart, Meder, Love & Speekenbrink